

Chapter 4. Binary Floating-Point Numbers

Machine representation of real numbers:

- Fixed-point (discussed in Chapter 1)
- Floating-point

Compared to fixed point numbers, floating-point numbers are efficient in representing both very large and very small numbers.

§1. IEEE 754-2008 Standard

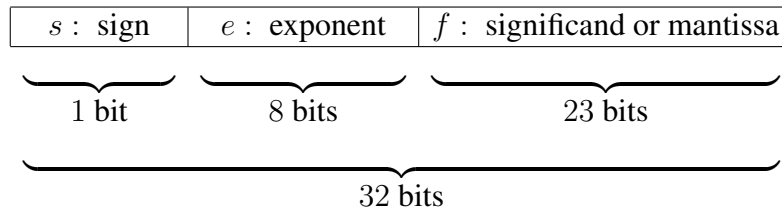
The IEEE 754-2008 standards are probably the most prevailing use of floating point numbers, which include four formats:

1. Single-precision format (32-bit)
2. Double-precision format (64-bit)
3. Single extended format (≥ 44 -bit)
4. Double extended format (≥ 80 -bit)

In this course we will introduce only the first two formats, IEEE single-precision format and IEEE double-precision format.

§1.1. IEEE single-precision (32 bits)

An IEEE single-precision floating-point number F has three parts, s , e and f , as shown below.



IEEE floating-point number F can be evaluated by

$$F = (-1)^s \times 1.f \times 2^{e-127}. \quad (1)$$

Note in (1) that there is a hidden 1 that is not shown in the representation of F . The maximal and minimal values of F can be decided by

$$\begin{aligned} F_{\max} &= (2 - 2^{-23}) \times 2^{254-127} = (1 - 2^{-24}) \times 2^{128}, \\ F_{\min} &= 1 \times 2^{1-127} = 2^{-126}. \end{aligned}$$

How to represent zero? Use the reserved formats as shown in the table below.

Value or meaning of an IEEE single-precision floating-point representation		
$e = 0$	$f = 0$	$F = \pm 0$
	$f \neq 0$	F are subnormal numbers ($= \pm 0.f \times 2^{-126}$)
$e = 255$	$f = 0$	$F = \pm \infty$
	$f \neq 0$	F is NAN (Not a Number)
$1 \leq e \leq 254$	—	F is an ordinary number and $F = (-1)^s \times 1.f \times 2^{e-127}$

Example 1 Convert 46.5_{10} into IEEE single-precision standard.

Solution: First we convert the given number 46.5_{10} into binary of the form of (1):

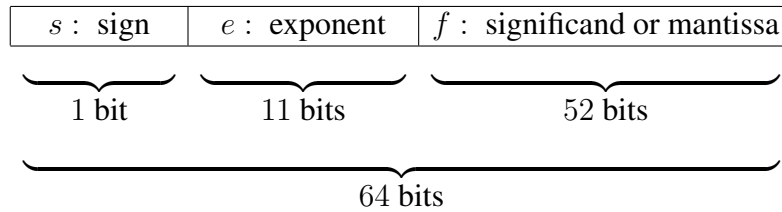
$$46.5_{10} = 101110.1_2 = 1.0111 \times 2^5 = (-1)^0 \times 1.0111 \times 2^{132-127}.$$

Then it is easy to obtain

$$\begin{cases} s &= 0 \\ e &= 1000\ 0100 \\ f &= 0111\ 0000\ 0000\ 0000\ 0000\ 000 \end{cases}$$

§1.2. IEEE double-precision (64 bits)

An IEEE double-precision floating-point number F has also three parts, s , e and f , as shown below.



IEEE floating-point number F can be evaluated by

$$F = (-1)^s \times 1.f \times 2^{e-1023}. \quad (2)$$

Note a hidden 1 also exists for double-precision case that is not shown in the representation of F . The maximal and minimal values of F can be decided by

$$\begin{aligned} F_{\max} &= (2 - 2^{-52}) \times 2^{2046-1023} = (1 - 2^{-53}) \times 2^{1024}, \\ F_{\min} &= 1 \times 2^{1-1023} = 2^{-1022}. \end{aligned}$$

The value or meaning of an IEEE double-precision number F is given in the following table:

Value or meaning of an IEEE double-precision floating-point representation		
$e = 0$	$f = 0$	$F = \pm 0$
	$f \neq 0$	F are subnormal numbers ($= \pm 0.f \times 2^{-1022}$)
$e = 1023$	$f = 0$	$F = \pm \infty$
	$f \neq 0$	F is NAN (Not a Number)
$1 \leq e \leq 1022$	—	F is an ordinary number and $F = (-1)^s \times 1.f \times 2^{e-1023}$

Example 2 Convert 46.5_{10} into IEEE double-precision standard.

Solution: First we convert the given number 46.5_{10} into binary of the form of (2):

$$46.5_{10} = 101110.1_2 = 1.0111 \times 2^5 = (-1)^0 \times 1.0111 \times 2^{1028-1023}.$$

Then it is easy to obtain

$$\begin{cases} s &= 0 \\ e &= 1000\ 0000\ 100 \\ f &= 0111\ 0000\ \underbrace{00 \dots 00}_{44\ \text{zeros}} \end{cases}$$

§2. A Summary of Floating-Point Representations

A floating-point representation F has three parts:

- The sign, S ;
- The significand (or mantissa), f (or M);
- The exponent, E .

The bits of the floating-point number F are stored in a register or a memory unit shown as follows

$$F = \boxed{S : \text{sign} \mid E : \text{exponent} \mid f \text{ or } M : \text{significand}}$$

and the value of such a floating-point number F is given by

$$F = (-1)^S \times M \times \beta^{E-\text{bias}}.$$

where β is the base of E which is implied in a system. Usually, β is often chosen as a power of 2.

The representation range is $[-F_{\max}, -F_{\min}]$ and $[F_{\min}, F_{\max}]$, where F_{\max} and F_{\min} are given by

$$\begin{aligned} F_{\max} &= f_{\max} \times \beta^{e_{\max}-\text{bias}}, \\ F_{\min} &= f_{\min} \times \beta^{e_{\min}-\text{bias}}. \end{aligned}$$

The computer system is overflow if a result is larger than F_{\max} or smaller than $-F_{\max}$. The system is called underflow if a result is nonzero and belongs to the interval $(-F_{\min}, F_{\min})$.

§3. Floating-Point Operations

The following discussions are using the generic floating point representation,

$$F = (-1)^S \times M \times \beta^{E-\text{Bias}}.$$

§3.1. Multiplication

Given two numbers

$$F_1 = (-1)^{S_1} \cdot M_1 \cdot \beta^{E_1 - \text{bias}} \text{ and } F_2 = (-1)^{S_2} \cdot M_2 \cdot \beta^{E_2 - \text{bias}},$$

and assume that they both are in normalized forms. Then the product $F_3 = F_1 \times F_2$ and

$$F_3 = (-1)^{S_3} \cdot M_3 \cdot \beta^{E_3 - \text{bias}}$$

can be obtained as follows.

1. Calculate $E_3 = E_1 + E_2 - \text{bias}$.
IF $(E_3 > E_{\max})$ THEN overflow ELSE IF $(E_3 < E_{\min})$ THEN underflow.
2. Calculate $M_3 = M_1 \times M_2$.
IF $(M_3 < (1/\beta))$ THEN $M_3 = M_3 \times \beta$ and $E_3 = E_3 - 1$.
3. IF $(E_3 < E_{\min})$ THEN underflow.

The second step above is called post-normalization.

§3.2. Addition/Subtraction

Let two float-point numbers be given by

$$\begin{aligned} F_1 &= (-1)^{S_1} \times M_1 \times \beta^{E_1 - \text{bias}} \\ F_2 &= (-1)^{S_2} \times M_2 \times \beta^{E_2 - \text{bias}} \end{aligned}$$

then addition/subtraction operation can be performed as follows.

1. Assume that $E_1 \geq E_2$ and compute

$$F_3 = F_1 \pm F_2 = [(-1)^{S_1} \times M_1 \pm (-1)^{S_2} \times M_2 \times \beta^{-(E_1 - E_2)}] \times \beta^{E_1 - \text{bias}}$$

2. Let $(-1)^{S_1} \times M_1 \pm (-1)^{S_2} \times M_2 \times \beta^{-(E_1 - E_2)}$ be denoted as M_3 . If $M_3 < (1/\beta)$ or $M_3 \geq 1$ then post-normalization is needed.

3. If there is post-normalization then we need to check whether or not the final exponent E_3 is overflow or underflow.

Example 3 Convert $F_1 = 4$ and $F_2 = 3$ into IEEE floating point single precision format. Then perform floating point operations $F_1 \times F_2$ and $F_1 + F_2$.

Solution:

$$F_1 = 4 = (-1)^{S_1} \times 1.f_1 \times 2^{E_1-127} = (-1)^0 \times 1.0 \times 2^{129-127},$$

$$F_2 = 3 = (-1)^{S_2} \times 1.f_2 \times 2^{E_2-127} = (-1)^0 \times 1.1 \times 2^{128-127}.$$

Then for floating point multiplication we have

$$F_3 = F_1 \times F_2 = (-1)^0 \times (1.0 \times 1.1) \times 2^{(129-127+128)-127} = (-1)^0 \times 1.1 \times 2^{130-127}.$$

For floating point addition, since $F_1 > F_2$ we first rewrite F_2 such that its 2's power part is the same as that of F_1 .

$$F_2 = (-1)^0 \times 1.1 \times 2^{128-127} = (-1)^0 \times 0.11 \times 2^{129-127}.$$

Then it follows

$$F_4 = F_1 + F_2 = (-1)^0 \times (1.0 + 0.11) \times 2^{129-127} = (-1)^0 \times 1.11 \times 2^{129-127}.$$

§4. Rounding Schemes

Rounding is a technique to obtain low-precision representation from given high-precision representation:

$$\text{High-precision} \rightarrow \text{Rounding} \rightarrow \text{Low-precision}$$

We assume that the input to a rounding scheme has m integer bits and d fractional bits,

$$x_{m-1} \dots x_0.x_{-1} \dots x_{-d},$$

and the output contains only integer bits, $y_{m'-1} \dots y_0$. This can be shown as

$$x_{m-1} \dots x_0.x_{-1} \dots x_{-d} \rightarrow \text{Rounding} \rightarrow y_{m'-1} \dots y_0.$$

Let X and $Y(X)$ denote the input and the output of a rounding scheme, respectively. Rounding error is defined as $R(X) = Y(X) - X$. We measure the accuracy of the rounding results by computing the maximum errors and the bias of the scheme, where bias is defined as the average error for a block of 2^d numbers including all the possible inputs to the rounding scheme.

Criteria for choosing a “good” rounding scheme (items with check mark will be discussed):

1. Accuracy of the output

- small maximum errors ✓
- small bias ✓
- small variation

2. Low time delay or speed ✓

3. Low implementation cost ✓

§4.1. Truncation or chopping: $\text{chop}(x)$

1. Definition: $\text{chop}(x) = \lfloor x \rfloor$.

2. Truth table:

Chopping scheme with $d = 2$		
Input: x	Output: $\text{chop}(x)$	Error: $\text{chop}(x) - x$
$\times.00$	$\times.$	0
$\times.01$	$\times.$	$-1/4$
$\times.10$	$\times.$	$-1/2$
$\times.11$	$\times.$	$-3/4$

3. Error and bias: It can be seen from the above table that the maximal error is $e_{\max}^- = -3/4$ and

$$\text{Bias} = \frac{1}{4} \left(0 - \frac{1}{4} - \frac{1}{2} - \frac{3}{4} \right) = -\frac{3}{8}.$$

4. Implementation: Its implementation is cost free.

5. Time delay: There is no delay incurred with this rounding scheme.

§4.2. Round to nearest integer: $\text{round}(x)$

1. definition: $\text{round}(x) = \lfloor x + 0.5 \rfloor$.

2. Truth table:

Round-to-nearest scheme with $d = 2$		
Input: x	Output: $\text{round}(x)$	Error: $\text{round}(x) - x$
$\times.00$	$\times.$	0
$\times.01$	$\times.$	$-1/4$
$\times.10$	$\times. + 1$	$+1/2$
$\times.11$	$\times. + 1$	$+1/4$

3. Error and bias: It can be seen from the above table that the maximal error is $e_{\max}^+ = +1/2$ and

$$\text{Bias} = \frac{1}{4} \left(0 - \frac{1}{4} + \frac{1}{2} + \frac{1}{4} \right) = +\frac{1}{8}.$$

4. Implementation: An implementation requires an adder and a few logic gates.

5. Time delay: It is equal to that of an adder of the size of the output.

§4.3. Round to nearest even integer: $\text{rtne}(x)$

1. definition: Round to the nearest even integer if it is a tie case.

2. Truth table:

Round-to-nearest-even scheme with $d = 2$		
Input: x	Output: $\text{rtne}(x)$	Error: $\text{rtne}(x) - x$
$\times 0.00$	$\times 0.$	0
$\times 0.01$	$\times 0.$	$-1/4$
$\times 0.10$	$\times 0.$	$-1/2$
$\times 0.11$	$\times 1.$	$+1/4$
$\times 1.00$	$\times 1.$	0
$\times 1.01$	$\times 1.$	$-1/4$
$\times 1.10$	$\times 1. + 1$	$+1/2$
$\times 1.11$	$\times 1. + 1$	$+1/4$

3. Errors and bias: It can be seen from the above table that the maximal errors are $e_{\max}^+ = +1/2$ and $e_{\max}^- = -1/2$.

$$\text{Bias} = \frac{1}{8} \left(0 - \frac{1}{4} - \frac{1}{2} + \frac{1}{4} + 0 - \frac{1}{4} + \frac{1}{2} + \frac{1}{4} \right) = 0.$$

4. Implementation and time delay: slightly higher than $\text{round}(x)$.

§4.4. ROM Rounding: $\text{ROM}(x)$

1. Definition: ROM rounding is given by

$$x_{m-1} \dots x_{\ell-1} x_{\ell-2} \dots x_0 . x_{-1} x_{-2} \dots x_{-d} \rightarrow \text{ROM}(x) \rightarrow x_{m-1} \dots x_{\ell-1} y_{\ell-2} \dots y_0.$$

Note that only ℓ bits are taken as actual input to ROM, $x_{\ell-2} \dots x_0 . x_{-1}$, and then ROM generates $\ell - 1$ output bits, $y_{\ell-2} \dots y_0$.

2. Truth table:

ROM scheme with $\ell = 3$ input bits		
Input: x	Output: $\text{ROM}(x)$	Error: $\text{ROM}(x) - x$
$\times 00.0$	$\times 00.$	0
$\times 00.1$	$\times 01.$	$+1/2$
$\times 01.0$	$\times 01.$	0
$\times 01.1$	$\times 10.$	$+1/2$
$\times 10.0$	$\times 10.$	0
$\times 10.1$	$\times 11.$	$+1/2$
$\times 11.0$	$\times 11.$	0
$\times 11.1$	$\times 11.$	$-1/2$

3. Error and bias: It can be seen from the above table that the maximal errors are $e_{\max}^+ = +1/2$ and $e_{\max}^- = -1/2$.

$$\text{Bias} = \frac{1}{8} \left(0 + \frac{1}{2} + 0 + \frac{1}{2} + 0 + \frac{1}{2} + 0 - \frac{1}{2} \right) = \frac{1}{8}.$$

4. Implementation cost: Size of the ROM used.

5. Time delay: Decided by the ROM reading speed.